

Apache Spark

2009: Conceived at UC Berkeley's AMPLab
Addressing MapReduce limitations

2010: Open-sourced
Gaining Traction

2013: Donated to Apache
Community-driven Open-Source

Present: Active Development
Large-scale Data Processing, ML

API Developed

PySpark (Python API)

Late 2010s: Emerged
User-Friendly Interface
Python for Distributed Computing

Present: Active Development
Large-scale Data Processing, ML

